

RAVISHEK KUMAR RANJAN

Bengaluru, India · ravishekr7@gmail.com · +91-6205194383 · LinkedIn · Portfolio

SUMMARY

Software Engineer with 4+ years of experience building scalable backend and data platforms on AWS. Strong in distributed systems, event-driven architectures, and large-scale data pipelines using Kafka, Spark, Iceberg, and S3. Experienced in asynchronous microservices and serverless workflows with Lambda, SQS, and Step Functions. Recently building AI-powered developer tooling and agent workflows using LangGraph, LangFlow, MCP servers, and LLM APIs.

SKILLS

Programming Languages:	Python, Java, C++, JavaScript, Scala
Backend & APIs:	REST APIs, Async APIs, Microservices, Event-driven Systems
Data & Streaming:	Apache Spark, Apache Kafka, Delta Lake, Apache Iceberg, ETL Pipelines
AI & Agents:	LangGraph, LangFlow, MCP (Model Context Protocol), OpenAI APIs, Prompt Engineering
Databases:	PostgreSQL, MySQL, MongoDB, DynamoDB
Cloud & Infra:	AWS (S3, EMR, Athena, Lambda, API Gateway, Step Functions, Glue, SQS, SNS, CloudWatch)
DevOps & Tools:	Docker, Git, CI/CD, Datadog, Linux
Certifications:	AWS Certified Data Engineer – Associate

WORK EXPERIENCE

Guidewire

Software Engineer 2

Bengaluru, India

Nov 2023 – Present

- Diagnosed and resolved **Iceberg small-file and metadata bloat issues** by building Spark-based compaction workflows in **AWS Glue**, reducing Athena query latency and lowering query costs across production datasets.
- Designed and implemented a **cloud-native distributed service** to detect and remediate data gaps across **event-driven data pipelines (Kafka, Iceberg, S3)**, ensuring end-to-end data consistency across multi-tenant environments.
- Led development of **asynchronous REST remediation APIs** using **API Gateway, Lambda, and Step Functions** to decode and republish lost Kafka messages across distributed services.
- Enhanced orchestration using **SQS + Lambda** to dynamically schedule DataGap validation jobs across **200+ environments**, introducing configurable concurrency controls and improving resource utilization.
- Built an Athena-based validation workflow to detect **PITR-induced data inconsistencies**, propagating invalidation signals across consumers to maintain post-recovery data integrity.

Tekion Corp

Associate Software Engineer

Bengaluru, India

June 2022 – Sept 2023

- Built scalable **Spark-based batch and streaming ETL pipelines** ingesting data from relational databases, Kafka, and Kinesis to produce fact, dimension, and aggregate datasets for analytics and ML workloads.
- Architected a **Delta Lake-based data warehouse on AWS S3**, enabling ACID guarantees, schema evolution, and high-performance query access via Athena, Presto, and EMR.
- Collaborated with analytics and ML teams to design curated datasets powering **financial reporting, dashboards, and model training pipelines**.
- Optimized ETL workflows by improving join strategies, DataFrame transformations, and S3 partitioning—reducing **pipeline runtime by 40%** and lowering **Athena query costs by 30%**.

PROJECTS

AI Email Automation Agent Pipeline (LangGraph, LangFlow, Python, MCP, OpenAI API, Gmail API)

- Designed a **multi-agent AI pipeline** using LangGraph and LangFlow to ingest and analyze Gmail messages.
- Built modular agents for **email classification, sender analysis, reporting, and automated unsubscribe workflows**, orchestrated through a stateful agent graph.
- Integrated **MCP (Model Context Protocol) servers** to enable structured tool access and contextual data retrieval for LLM agents.
- Implemented LLM-driven analysis to identify recurring spam senders and generate structured summaries using **agent-based AI workflows**.

EDUCATION

Dayananda Sagar College of Engineering

BE, Information Science Engineering

Bengaluru, India

Aug 2018 – June 2022